



TECHNICAL UNIVERSITY OF DENMARK

Trading ETFs

Student :

VISNUKARAN KIRUBAKARAN,
STUDENT ID: s224527

Teacher :

October 27, 2024

Contents

1	Descriptive analysis	2
	a) Description of the data material	2
	b) Density histogram of the weekly returns from the ETF AGG	2
	c) The weekly return over time for each of the four ETFs	3
	d) Box plot of the weekly returns by ETF	5
	e) Summary sizes for the four ETFs	6
2	Statistical analysis	7
	f) Statistical models describing the weekly return for each of the four ETFs	7
	g) Confidence intervals	9
	h) Hypothesis test	9
	i) Hypothesis Test: Comparison of Weekly Returns from VAW and AGG .	10
	j) The importance of statistical test	11
	k) Correlation between ETFs	11

1 Descriptive analysis

a) Description of the data material

The dataset consists of weekly returns for 95 ETFs. It contains 96 columns, with the first column representing the date, and the remaining columns showing the weekly returns for each ETF. In this project, the focus will be on 4 specific ETFs: AGG, VAW, IWN, and SPY. These are considered quantitative variables, as they represent numerical data that can be analyzed and computed. For each ETF, there are a total of 454 observations, starting from May 5, 2006 and ending on May 8, 2015. However, it is noted that there are some missing values in the dataset. This is evident from the irregular intervals between certain dates, indicating that data for some weeks is missing.

b) Density histogram of the weekly returns from the ETF AGG

Below can the density histogram of the weekly returns from the ETF AGG be seen with an overlaid normal curve.




Figure 1: Density histogram of the weekly returns from ETF AGG with normal distribution overlay.

The histogram can be seen as symmetrical because the mean and median of the ETF returns are close, indicating a balanced distribution of values around the center. This lack of skewness, combined with the equal spread of data on both sides of the mean, suggests the returns follow a roughly normal distribution. Furthermore, the return can both be seen in the positive and negative directions of the x-axis, meaning the return can

be both. Lastly, the histogram moderate variation in the data, as the data points are spread around the mean with a roughly bell-shaped curve. The central peak indicates that most observations are clustered around the mean, while the tails show fewer extreme values.

c) The weekly return over time for each of the four ETFs

The four plots illustrate the weekly returns of different ETFs (AGG, VAW, IWN, and SPY) over time, showing distinct volatility patterns for each. AGG, being a bond ETF, has much smaller fluctuations in returns, reflecting its lower risk profile compared to the others. In contrast, VAW and IWN display larger and more frequent swings, indicating higher potential for both gains and losses. A common feature across all ETFs is the period of significant volatility between 2008 and 2009, likely due to the global financial crisis. Some instability is also observed around 2011, reflecting continued market uncertainty. These fluctuations suggest varying levels of market sensitivity, with AGG being more stable and VAW/IWN showing greater exposure to risk.




Figure 2: Figure 2: AGG's development over time




Figure 3: Figure 3: VAW's development over time




Figure 4: Figure 4: IWN's development over time




Figure 5: Figure 5: SPY's development over time

d) Box plot of the weekly returns by ETF

The box plot illustrates the empirical distribution of the weekly returns from each of the four ETFs




Figure 6: Boxplot of the weekly returns by ETF.

The box plot shows the distribution of weekly returns for the four ETFs (AGG, VAW, IWN, and SPY). AGG has the most compact and nearly symmetrical distribution, while VAW and SPY are left-skewed, with more negative returns. IWN is the most balanced, with returns spread symmetrically around the median. VAW has the largest variability, shown by its wide interquartile range (IQR), while AGG has the smallest. All ETFs

display several outliers, particularly VAW and IWN, which could be due to market events during the period.

e) Summary sizes for the four ETFs

	Number of obs.	Sample mean	Sample variance	Std. dev.	Lower quartile (Q1)	Median (Q2)	Upper quartile (Q3)
AGG	454.000000	0.000266	0.000036	0.005976	- 0.002973	0.000237	0.003893
VAW	454.000000	0.001794	0.001302	0.036083	- 0.016096	0.004798	0.019685
IWN	454.000000	0.001188	0.001025	0.032015	- 0.014305	0.003120	0.019056
SPY	454.000000	0.001360	0.000614	0.024786	- 0.011325	0.004216	0.014498

Table 1: Summary sizes for the four ETFs

The table provides precise values for the mean, variance, and standard deviation, which are not visible in the box plot. It also shows the exact number of observations and quartiles, offering more detailed insights than the visual representation of the box plot.

2 Statistical analysis

f) Statistical models describing the weekly return for each of the four ETFs

The weekly returns from the four ETF's can be separated into the following statistical models:

ETF	Statistical Model
AGG	$AGG_i \sim \mathcal{N}(\mu_{AGG}, \sigma_{AGG}^2)$ and i.i.d., where $i = 1, \dots, 454$
VAW	$VAW_i \sim \mathcal{N}(\mu_{VAW}, \sigma_{VAW}^2)$ and i.i.d., where $i = 1, \dots, 454$
IWN	$IWN_i \sim \mathcal{N}(\mu_{IWN}, \sigma_{IWN}^2)$ and i.i.d., where $i = 1, \dots, 454$
SPY	$SPY_i \sim \mathcal{N}(\mu_{SPY}, \sigma_{SPY}^2)$ and i.i.d., where $i = 1, \dots, 454$

Table 2: Table 2: Statistical models for the weekly returns of the four ETFs

It is assumed that the observations in these statistical models are independent of each other, have the same distribution, and follow a normal distribution.

The parameters for the models can furthermore be estimated from the sample mean and the sample variances. These parameters were already found in section 1, e, and can be seen here:

	Sample mean	Sample variance
AGG	0.000266	0.000036
VAW	0.001794	0.001302
IWN	0.001188	0.001025
SPY	0.001360	0.000614

Table 3: Estimated parameters

To perform model validation, QQ-plots for the four ETFs (AGG, VAW, IWN, and SPY) were compared with QQ-plots from simulated normal data.

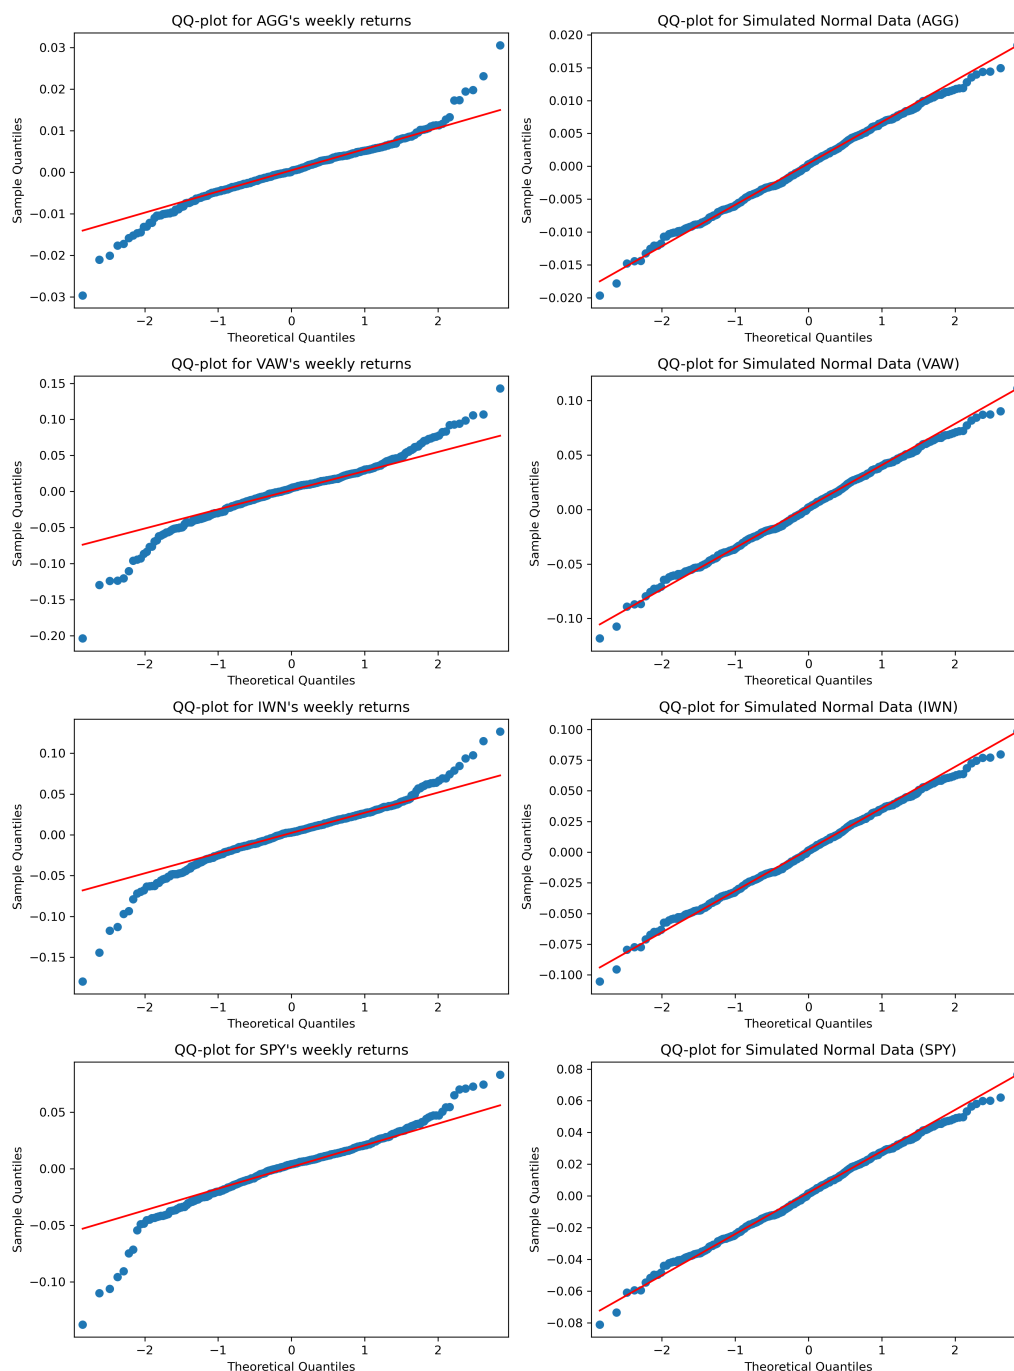


Figure 7: QQ-plots of the ETF's and simulations.

The observed ETF data deviates from normality, particularly in the tails, which suggests some degree of non-normality in the weekly returns. In contrast, the simulated data follows the expected normal distribution more closely. However, due to the large sample size (454 observations for each ETF), the Central Limit Theorem (CLT) becomes important. The CLT states that the sample mean will approximate a normal distribution,

even if the underlying data is not perfectly normal. Thus, the models are still reliable for further analysis, supported by the CLT, which can be expressed as:

$$\frac{\bar{X} - \mu}{\sigma/\sqrt{n}} \sim N(0, 1^2) \quad (1)$$

g) Confidence intervals

To calculate the 95% confidence interval (CI) for the mean weekly return of AGG, we use the t-distribution with n-1 degrees of freedom, where n is the sample size (454 observations). The 0.975 quantile from the t-distribution is used to capture the central 95% of the distribution. The following equation can be used:

$$\bar{x} \pm t_{1-\alpha/2} \cdot \frac{s}{\sqrt{n}}$$

We can find the critical t value using simple python where we get the value 1.9652145681681557. Now we can calculate the CI for the mean weekly return of AGG using the formula: The mean we found in section f) to be 0.000265757 and the sample standard deviation from section e) to be 0.00597584. Now we can calculate it for AGG. AGG: $\bar{x} \pm t_{1-\alpha/2} \cdot \frac{s}{\sqrt{n}} = 0.000265757 \pm 1.97 \cdot \frac{0.005976}{\sqrt{454}} = -0.0002867641$ and 0.0008182781 . We can compute the corresponding intervals for the three remaining ETFs and fill it in a table:

ETF	Lower bound of CI	Upper bound of CI
AGG	-0.000285407345518444	0.0008169212976295796
VAW	-0.0015342078012798403	0.005121788040169244
IWN	-0.0017651741507074973	0.004140532646420902
SPY	-0.0009259600489804174	0.0036461709486724178

Table 4: Confidence intervals for the four ETFs

When we compare the calculated CI for AGG with the given python code, it can be seen that the results are nearly equal with small rounding differences.

h) Hypothesis test

To test whether the mean weekly return for AGG deviates significantly from zero, we set up the following hypotheses:

$$H_0 : \mu_{AGG} = 0 \quad \text{and} \quad H_1 : \mu_{AGG} \neq 0$$

We perform the test at the 5% significance level ($\alpha = 0.05$) and calculate the t-statistic using the formula:

$$t = \frac{\bar{x} - \mu_0}{s/\sqrt{n}}$$

We can insert the values, and then get:

$$t_{obs} = \frac{0.000265757 - 0}{\frac{0.005975841}{\sqrt{454}}} = 0.9475750245$$

The p value can be calculated as follows:

$$p\text{-value} = 2 \cdot P(T > |t_{obs}|) = 2 \cdot P(1.97 > 0.948) = 0.3439$$

Conclusion: The hypothesis test suggests that the mean weekly return for AGG does not significantly deviate from zero. This conclusion aligns with the confidence interval calculated earlier, which also contained zero within its range, further supporting the result that the mean return is not statistically different from zero at the 95% confidence level. Thus, performing the hypothesis test confirms the conclusion we obtained from the confidence interval, and we do not need to reject the model. Furthermore, the calculated results are equal to the results found in python.

i) Hypothesis Test: Comparison of Weekly Returns from VAW and AGG

To investigate whether the mean weekly returns of VAW and AGG differ, we perform a Welch two-sample t-test. The hypothesis test is formulated as follows:

$$H_0 : \mu_{VAW} = \mu_{AGG}$$

$$H_1 : \mu_{VAW} \neq \mu_{AGG}$$

The null hypothesis H_0 states that the mean weekly returns of VAW and AGG are equal, while the alternative hypothesis H_1 states that the mean weekly returns differ.

Significance Level:

We set the significance level at $\alpha = 0.05$.

Test Statistic:

Using the Welch two-sample t-test, the test statistic is calculated as:

$$t_{obs} = \frac{(\bar{x}_{VAW} - \bar{x}_{AGG}) - \delta_0}{\sqrt{\frac{s_{VAW}^2}{n_{VAW}} + \frac{s_{AGG}^2}{n_{AGG}}}}$$

where:

- $\bar{x}_{VAW} = 0.001794$ and $\bar{x}_{AGG} = 0.000266$ are the sample means of VAW and AGG.
- $s_{VAW}^2 = 0.001302$ and $s_{AGG}^2 = 0.000036$ are the sample variances.
- $n_{VAW} = 454$ and $n_{AGG} = 454$ are the sample sizes for VAW and AGG.

If we insert the values we get:

$$t_{\text{obs}} = \frac{(0.000265757 - 0.00179379) - 0}{\sqrt{\frac{0.005975841^2}{454} + \frac{0.03608286^2}{454}}} = -0.890$$

The degrees of freedom for the Welch t-test are calculated using the following formula:

$$\nu = \frac{\left(\frac{s_{\text{VAW}}^2}{n_{\text{VAW}}} + \frac{s_{\text{AGG}}^2}{n_{\text{AGG}}}\right)^2}{\frac{\left(\frac{s_{\text{VAW}}^2}{n_{\text{VAW}}}\right)^2}{n_{\text{VAW}}-1} + \frac{\left(\frac{s_{\text{AGG}}^2}{n_{\text{AGG}}}\right)^2}{n_{\text{AGG}}-1}}$$

We again insert the values and get:

$$v = \frac{\left(\frac{0.005975841^2}{454} + \frac{0.03608286^2}{454}\right)^2}{\frac{\left(\frac{0.005975841^2}{454}\right)^2}{454-1} + \frac{\left(\frac{0.03608286^2}{454}\right)^2}{454-1}} = 478$$

The test statistic is computed as 0.890:

The corresponding p-value is 0.374.

Since the p-value of 0.374 is greater than the significance level of 0.05, we fail to reject the null hypothesis. This indicates that there is insufficient evidence to conclude that the mean weekly returns of VAW and AGG differ significantly.

Thus, we conclude that the mean weekly returns of VAW and AGG are not significantly different, and there is no statistical evidence to suggest that one ETF has a higher mean return than the other. Furthermore, when we compare it to the python result, we get exact the same answer.

j) The importance of statistical test

In task i, conducting the hypothesis test was essential because the confidence intervals from both ETFs overlapped. This overlap means the confidence intervals alone couldn't provide a clear conclusion about whether the mean returns differed. When confidence intervals overlap, additional testing is needed to draw any significant conclusions. Had there been no overlap, we could have dismissed the null hypothesis using just the confidence intervals. But since there was overlap, both the confidence intervals and the hypothesis test with a significance level of $\alpha = 5\%$ were required to arrive at a reliable result.

k) Correlation between ETFs

When constructing a portfolio of ETFs, the diversification of risk is crucial to reduce the overall volatility of the portfolio. One way to achieve diversification is by combining assets with low or negative correlations. The correlation between two ETFs measures how the

returns of one ETF move relative to the other. In general, a negative or low correlation between two ETFs indicates a higher potential for diversification, as one ETF's gain may offset the other's loss, reducing the overall portfolio risk.

The sample covariance between two ETFs is computed using the following formula:

$$s_{xy} = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})$$

where:

- s_{xy} is the sample covariance,
- x_i and y_i are the individual returns for the two ETFs,
- \bar{x} and \bar{y} are the means of the returns for each ETF,
- n is the number of observations.

The sample correlation coefficient is calculated as follows:

$$r = \frac{s_{xy}}{s_x \cdot s_y}$$

where:

- s_x and s_y are the sample standard deviations of the two ETFs,
- s_{xy} is the sample covariance computed above.

We can insert our values, and calculate the correlation:

$$r = \frac{0.000984}{0.036083 \cdot 0.032015} = 0.852$$

This high positive correlation suggests that the weekly returns of VAW and IWN tend to move together in the same direction. This is consistent with the expectation, as both ETFs may be influenced by similar market factors.

We can further illustrate this relationship by creating a scatter plot of the weekly returns of VAW against IWN.




Figure 8: Scatter plot of weekly returns for VAW vs. IWN.

The scatter plot above demonstrates a positive linear relationship between the returns of VAW and IWN, which aligns with the high correlation coefficient observed. A strong positive correlation implies that the returns of both ETFs move similarly, indicating less opportunity for risk diversification by combining them in a portfolio.

In conclusion, while combining VAW and IWN might still be part of a larger portfolio strategy, they would not provide substantial diversification benefits due to their high correlation. If we compare the correlation computed above to the python code, we get the exact same result.